

## 1 Pre-Check

This section is designed as a conceptual check for you to determine if you conceptually understand and have any misconceptions about this topic. Please answer true/false to the following questions, and include an explanation:

- 1.1 MapReduce is more general than Spark since it is lower level.

false.

→ specific  
high level → abstract

- 1.2 The higher the PUE the more efficient the datacenter is.

power usage efficiency

- 1.3 Hamming codes can detect any type of data corruption.

false

- 1.4 All RAID levels improve reliability.

false.

## 2 Hamming ECC

Recall the basic structure of a Hamming code. We start out with some bitstring, and then add parity bits at the indices that are powers of two (1, 2, 8, etc.). We don't assign values to these parity bits yet. **Note that the indexing convention used for Hamming ECC is different from what you are familiar with.** In particular, the 1 index represents the MSB, and we index from left-to-right. The  $i$ th parity bit  $P\{i\}$  covers the bits in the new bitstring where the *index* of the bit under the aforementioned convention,  $j$ , has a 1 at the same position as  $i$  when represented as binary. For instance, 4 is 0b100 in binary. The integers  $j$  that have a 1 in the same position when represented in binary are 4, 5, 6, 7, 12, 13, etc. Therefore,  $P_4$  covers the bits at indices 4, 5, 6, 7, 12, 13, etc. A visual representation of this is:

Bit position	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	
Encoded data bits	p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11	p16	d12	d13	d14	d15	
Parity bit coverage	p1	x		x		x		x		x		x		x		x		x		x	
	p2		x	x					x	x				x	x				x	x	
	p4				x	x	x	x					x	x	x	x					x
	p8								x	x	x	x	x	x	x						
	p16																x	x	x	x	x

Source: [https://en.wikipedia.org/wiki/Hamming\\_code](https://en.wikipedia.org/wiki/Hamming_code)

- 2.1 How many bits do we need to add to 0011<sub>2</sub> to allow single error correction?

1 2 3 4 5 6 7  
x x 0 x 0 1 1

4

0000011

2.2 Which locations in  $0011_2$  would parity bits be included?

2.3 Which bits does each parity bit cover in  $0011_2$ ?

$P_1$  1 3 5 7  
 $P_2$  2 3 6 7  
 $P_4$  4 5 6 7

2.4 Write the completed coded representation for  $0011_2$  to enable single error correction. Assume that we set the parity bits so that the bits they cover have even parity.

2.5 How can we enable an additional double error detection on top of this?

Additional bit

2.6 Find the original bits given the following SEC Hamming Code:  $011011_2$ . Again, assume that the parity bits are set so that the bits they cover have even parity.

$1\ 2\ 3\ 4\ 5\ 6\ 7$   
 $0\ 1\ 1\ 0\ 1\ 1$   
 $1\ 1\ 1\ 1\ 1$   
 $0\ 1\ 1\ 0\ 0\ 1\ 1$   
 $1\ 1\ 1\ 1$   
 $2\ 3\ 4\ 5\ 6\ 7$

$P_1 \rightarrow e.$   
 $P_2$   
 $P_4 \rightarrow e$

2.7 Find the original bits given the following SEC Hamming Code:  $1001000_2$

$P_1 \rightarrow e$   
 $P_2$   
 $P_4 \rightarrow e.$

1001000

### 3 RAID

3.1 Fill out the following table:

	Configuration	Pro/Good for	Con/Bad for
RAID 0	data multiple disk	no overhead fast r/w	reliability
RAID 1	mirrored disk	fast r/w fast recovery	high overhead expensive.
RAID 2	Hamming ecc. bit level striping	smaller overhead	redundant check disk
RAID 3	bit level striping	smallest overhead	need all disk
RAID 4	block level striping	higher throughput for small read	slow small writes

RAID 5	work local striping parity distributed across disk	higher throughput of small use	keeps 1 line very
--------	--	--------------------------------	-------------------

## 4 MapReduce

For each problem below, write pseudocode to complete the implementations. Tips:

- The input to each MapReduce job is given by the signature of `map()`.
- `emit(key k, value v)` outputs the key-value pair `(k, v)`.
- `for var in list` can be used to iterate through `Iterables` or you can call the `hasNext()` and `next()` functions.
- Usable data types: `int`, `float`, `String`. You may also use lists and custom data types composed of the aforementioned types.
- `intersection(list1, list2)` returns a list of the common elements of `list1`, `list2`.

- 4.1 Given a set of coins and each coin's owner in the form of a list of `CoinPairs`, compute the number of coins of each denomination that a person has

CoinPair:  
String person  
String coinType

1 map(CoinPair pair):

*emit(pair, 1)*

1 reduce(\_\_\_\_\_, \_\_\_\_\_):

*coin pair pair Iterable(CoinPair)*

*total = 0  
for num in count:  
total += num  
emit(pair, total)*

- 4.2 Using the output of the first MapReduce, compute each person's amount of money.  
`valueOfCoin(String coinType)` returns a float corresponding to the dollar value of the coin.

1 map(tuple<CoinPair, int> output):

*emit((pair, coinPair, person),  
valueOfCoin(pair.coinType) \* count)*

1 reduce(\_\_\_\_\_, \_\_\_\_\_):

*coin pair, person Iterable(CoinPair)*

*total = 0  
for val in value:  
total += val*

*emit(person, -total)*

## 5 Spark

**Resilient Distributed Datasets (RDD)** are the primary abstraction of a distributed collection of items

**Transforms**  $RDD \rightarrow RDD$

**map( $f$ )** Return a new transformed item formed by calling  $f$  on a source element.

**flatMap( $f$ )** Similar to map, but each input item can be mapped to 0 or more output items (so  $f$  should return a sequence rather than a single item).

**reduceByKey( $f$ )** When called on a dataset of  $(K, V)$  pairs, returns a dataset of  $(K, V)$  pairs where the values for each key are aggregated using the given reduce function  $f$ , which must be of type  $(V, V) \rightarrow V$ .

**Actions**  $RDD \rightarrow Value$

**reduce( $f$ )** Aggregate the elements of the dataset regardless of keys using a function  $f$ .

Call `sc.parallelize(data)` to parallelize a Python collection, data.

- 5.1 Given a set of coins and each coin's owner, compute the number of coins of each denomination that a person has. Then, using the output of the first result, compute each person's amount of money. Assume `valueOfCoin(coinType)` is defined and returns the dollar value of the coin.

The type of `coinPairs` is a tuple of (person, coinType) pairs.

```
1 coinData = sc.parallelize(coinPairs)
  Out1 = coinData.map(lambda (k,v): (k,v[1]))
              ~ reduce by key (lambda (v1,v2): v1+v2)
  Out2 = Out1.map(lambda (k,v): (k[0], v * values[coinType[i]]))
              ~ reduce by key (lambda v1, v2: v1+v2)
```

- 5.2 Given a student's name and course taken, output their name and total GPA.

CourseData:

`int` courseID

`float` studentGrade // a number from 0-4

The type of `students` is a list of (studentName, courseData) pairs.

```
1 studentsData = sc.parallelize(students)
```

Out = studentsData.map(lambda (k,v): (k, (v[1].studentGrade, 1)))

reduce by key (lambda (v1,v2): [v1[0]+v2[0], v1[1]+v2[1]])

map (lambda (k,v): (k[0], v[1]))

Out = studentsData.map(lambda (k,v): (k, v[1]))

reduce by key (lambda (v1,v2): (v1+v2))

map (lambda (k,v): ((k, v), k[1]))

reduce by key (lambda (k,v): (k, v/k[1]))

computer room air conditioner

## 6 Warehouse-Scale Computing

Sources speculate Google has over 1 million servers. Assume each of the 1 million servers draw an average of 200W, the PUE is 1.5, and that Google pays an average of 6 cents per kilowatt-hour for datacenter electricity.

- 6.1 Estimate Google's annual power bill for its datacenters.


$$1.5 \times 10^6 \times 0.2 \text{ kW} \times 0.06 \times 8760 \text{ h/y}$$

- 6.2 Google reduced the PUE of a 50,000-machine datacenter from 1.5 to 1.25 without decreasing the power supplied to the servers. What's the cost savings per year?

$$\text{PUE} = \frac{\text{Total building power}}{\text{IT equipment power}}$$

$$50,000 \times (1.5 - 1.25) \times 0.2 \text{ kW} \times 0.06 \times 8760 \text{ h/y}$$

Servers + networking



CRAAC  
Lighting  
chiller  
UPS  
PDU  
transformers  
switches